

Validation

December 7, 2020

```
[ ]: ![ -e 'zip.train' ] || ( wget https://web.stanford.edu/~hastie/ElemStatLearn/  
    ↳ datasets/zip.train.gz && gzip -d zip.train.gz )  
![ -e 'zip.test' ] || ( wget https://web.stanford.edu/~hastie/ElemStatLearn/  
    ↳ datasets/zip.test.gz && gzip -d zip.test.gz )
```

Import necessary packages: Numpy, Pandas, matplotlib

```
[ ]: import numpy as np  
import pandas as pd  
from matplotlib import pyplot as plt
```

Mount your google drive (if you have a google account) or upload files (go on the file icon on the left -> right click). Copy path of zip.train and zip.test and load them as numpy arrays using the following code (insert the path as string).

```
[ ]: #dirname = '/content/drive/My Drive/ML_Class_2020/KNN/'  
dirname = './'  
path_to_train = dirname + 'zip.train'  
path_to_test = dirname + 'zip.test'  
training_data = np.array(pd.read_csv(path_to_train, sep=' ', header=None))  
test_data = np.array(pd.read_csv(path_to_test, sep=' ', header=None))  
  
X_train, y_train = training_data[:,1:-1], training_data[:,0]  
X_test, y_test = test_data[:,1:], test_data[:,0]  
  
# We only want to classify two different digits. Note the only difference is  
# that we don't load as many ones as we load zeros.  
  
X_train_0 = X_train[y_train == 0]  
X_train_1 = X_train[y_train == 1][:30]  
X_train = np.vstack((X_train_0, X_train_1))  
  
y_train_0 = y_train[y_train == 0]  
y_train_1 = y_train[y_train == 1][:30]  
y_train = np.hstack((y_train_0, y_train_1))  
  
X_test_0 = X_test[y_test == 0]  
X_test_1 = X_test[y_test == 1][:12]
```

```
X_test = np.vstack((X_test_0, X_test_1))

y_test_0 = y_test[y_test == 0]
y_test_1 = y_test[y_test == 1][:12]
y_test = np.hstack((y_test_0, y_test_1))
print(y_test.shape)
```

(371,)


```
[ ]: def show_numbers(X):
    num_samples = 90
    indices = np.random.choice(range(len(X)), num_samples)
    print(indices.shape)
    sample_digits = X[indices]

    fig = plt.figure(figsize=(20, 6))

    for i in range(num_samples):
        ax = plt.subplot(6, 15, i + 1)
        img = 1-sample_digits[i].reshape((16, 16))
        plt.imshow(img, cmap='gray')
        plt.axis('off')
```

```
[ ]: show_numbers(X_test)
```

(90,)



Choose one of the classifiers you implemented and use it on the following task: Take a look at the provided notebook. It loads the ZIP-Code's zeros and ones, but now there are many more zeros than there are ones. We are going to believe that these digits represent genetic data. Actually, we are virologists trying to develop a test to determine which of the data belongs to infected patients (the ones) and which do not (the zeros). Train/apply your model on the data and do the tasks of the exercise.

```
[ ]: from collections import Counter
from functools import wraps, partial

def collect_after(c):
    def decorator(f):
        @wraps(f)
        def g(*args, **kwargs):
            return c(f(*args, **kwargs))
        return g
    return decorator

class KNearestNeighbors():
    """
    Think about defining more functions that will help you building this
    ↪ algorithm.
    Optimally, one that takes in k and a test image as a parameter.
    """
    def squared_euclidean_distance(self, x_1, x_2, axis=1):
        """
        np.sum(x, axis = 1) will be summing all elements over the pixel dimension
        ↪ (axis = 1)
        """
        return np.sum((x_1 - x_2) ** 2, axis=axis)

    def set_k(self, k):
        self.k = k

    def __init__(self, k):
        self.set_k(k)

    def fit(self, X, y):
        self.X = X
        self.y = y

    @collect_after(tuple)
    def calculate_nearest(self, X, k=None):
        if k is None:
            k = self.k
        for obj in X:
            nearest = []
            for i, known in enumerate(self.X):
                d = self.squared_euclidean_distance(known, obj, axis=0)
                nearest.append((d, i))
            nearest.sort()
            nearest[k:] = []
```

```

        yield tuple(nearest)

    @collect_after(partial(np.fromiter, dtype=float))
    def predict_class(self, X, nearest=None):
        res = []
        k = self.k
        if nearest is None:
            nearest = self.calculate_nearest(X, k)
        for nearest_current in nearest:
            yield Counter(self.y[i] for _, i in nearest_current[:k]).
↳most_common(1)[0][0]

    def predict_class_different_k(self, X, ks=(3,)):
        ks = tuple(ks)
        nearest = self.calculate_nearest(X, max(ks))
        for k in ks:
            self.set_k(k)
            yield (k, self.predict_class(X, nearest=nearest))

def accuracy(y_actual, y_pred):
    return (y_actual == y_pred).sum() / len(y_actual)

```

```

[ ]: @collect_after(tuple)
def k_fold_split(X, y, k):
    n = len(y)
    permut = np.argsort(np.random.random_sample((n,)))
    X_s = X[permut]
    y_s = y[permut]
    for i in range(k):
        start, end = int(i * n / k), int((i + 1) * n / k)
        yield (X_s[start:end], y_s[start:end])

```

```

[ ]: class CrossValidation():
    def stack_pairs(self, pairs):
        return np.vstack([Xy[0] for Xy in pairs]), np.hstack([Xy[1] for Xy in
↳pairs])

    def joins(self, samples):
        for i in range(len(samples)):
            yield (samples[i], self.stack_pairs(samples[:i] + samples[i + 1:]))

    def __init__(self, model_class, samples, params, loss):
        self.models = []
        self.train_val_splits = []
        losses = []
        for Xy_val, Xy_train in self.joins(samples):

```

```

        self.train_val_splits.append((Xy_train, Xy_val))
        model = model_class(**params)
        self.models.append(model)
        model.fit(*Xy_train)
        y_pred = model.predict_class(Xy_val[0])
        L = loss(Xy_val[1], y_pred)
        losses.append(L)
    self.losses = np.array(losses)
    self.bias = np.mean(self.losses)
    self.var = np.var(self.losses)

```

```

[ ]: def one_minus_accuracy(y_actual, y_pred):
    return 1.0 - accuracy(y_actual, y_pred)

```

```

[ ]: from itertools import product
def iter_params(spaces):
    names = spaces.keys()
    for values in product(*(spaces[k] for k in names)):
        yield dict(zip(names, values))

```

```

[ ]: def k_fold_search_params(model_class, X, y, k, spaces, loss):
    info = {**{'param_' + i: [] for i in spaces.keys()}, 'bias': [], 'var': []}
    splits = k_fold_split(X, y, k)
    for params in iter_params(spaces):
        cv = CrossValidation(model_class, splits, params, loss)
        info['bias'].append(cv.bias)
        info['var'].append(cv.var)
        for p, v in params.items():
            info['param_' + p].append(v)
    return pd.DataFrame(info)

```

```

[ ]: %%time
np.random.seed(42)
search_df = k_fold_search_params(KNearestNeighbors, X_train, y_train, 5, {'k':
→np.arange(1, 30, 1, dtype=int)}, one_minus_accuracy)

```

CPU times: user 5min 6s, sys: 57.8 ms, total: 5min 6s
 Wall time: 5min 7s

```

[ ]: search_df

```


```

[ ]:
   param_k  bias  var
0         1  0.000000  0.000000
1         2  0.000000  0.000000
2         3  0.000000  0.000000
3         4  0.000000  0.000000
4         5  0.000000  0.000000

```

5	6	0.000000	0.000000
6	7	0.000000	0.000000
7	8	0.000000	0.000000
8	9	0.000000	0.000000
9	10	0.000000	0.000000
10	11	0.000000	0.000000
11	12	0.000000	0.000000
12	13	0.000000	0.000000
13	14	0.000000	0.000000
14	15	0.000000	0.000000
15	16	0.000000	0.000000
16	17	0.000000	0.000000
17	18	0.000000	0.000000
18	19	0.000000	0.000000
19	20	0.000000	0.000000
20	21	0.000000	0.000000
21	22	0.000000	0.000000
22	23	0.000000	0.000000
23	24	0.000000	0.000000
24	25	0.000816	0.000003
25	26	0.000000	0.000000
26	27	0.000816	0.000003
27	28	0.000816	0.000003
28	29	0.000816	0.000003

```
[ ]: plt.figure(figsize=(10, 6))
plt.plot(search_df['param_k'], search_df['bias'], linestyle='--',
        color='#FF0000', label='bias')
plt.plot(search_df['param_k'], search_df['var'], linestyle=':',
        color='#00FF00', label='var')
plt.xlabel('k')
plt.legend()
plt.show()
```



The model seems to almost always give the correct answer on the training dataset.

```
[ ]: from collections import namedtuple
from functools import lru_cache

class Confusion(namedtuple('Confusion', ['TP', 'FP', 'FN', 'TN'])):
    def calculate(y_true, y_pred):
        y_true = y_true == True
        y_pred = y_pred == True
        return Confusion((y_true & y_pred).sum(), (~y_true & y_pred).sum(),
                          (y_true & ~y_pred).sum(), (~y_true & ~y_pred).sum())

    @property
    def P(self):
        return self.TP + self.FN

    @property
    def N(self):
        return self.FP + self.TN

    @property
    def PP(self):
        return self.TP + self.FP

    @property
```

```

def PN(self):
    return self.FN + self.TN

# Precision
@property
@lru_cache()
def PPV(self):
    return self.TP / self.PP

# Recall
@property
@lru_cache()
def TPR(self):
    return self.TP / self.P

@property
def F1(self):
    return 2 / (1 / self.TPR + 1 / self.PPV)

@property
def accuracy(self):
    return (self.TP + self.TN) / (self.P + self.N)

```

```

[ ]: %%time
model = KNearestNeighbors(k=3)
model.fit(X_train, y_train)
conf = Confusion.calculate(y_test, model.predict_class(X_test))

```

CPU times: user 3.6 s, sys: 998 µs, total: 3.6 s
Wall time: 3.6 s

```
[ ]: conf
```

```
[ ]: Confusion(TP=12, FP=1, FN=0, TN=358)
```

```

[ ]: def print_scores(conf):
    print(f"""
Precision: {conf.PPV:.3}
Recall: {conf.TPR:.3}
F1-score: {conf.F1:.3}
Accuracy: {conf.accuracy:.3}
""")

```

```
[ ]: print_scores(conf)
```

Precision: 0.923

Recall: 1.0
F1-score: 0.96
Accuracy: 0.997

```
[ ]: conf_zeros = Confusion.calculate(y_test, np.zeros_like(y_test))  
     print_scores(conf_zeros)
```

Precision: nan
Recall: 0.0
F1-score: nan
Accuracy: 0.968

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:31: RuntimeWarning:  
invalid value encountered in long_scalars  
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:41: RuntimeWarning:  
divide by zero encountered in double_scalars
```

Precision is not defined, because we have never classified objects as positive, but F1 would be 0 in similar cases.

Precision only says the proportion of correct answers. F1-score shows, how the model actually is at distinguishing the two classes.

According to the given test dataset, the probability is 92. Yes, I would.

But even if the probabilities had been less, using the test could be recommended depending on the cost of the two types of error,