

Perceptron

January 18, 2021

1 Mustererkennung/Machine Learning - Assignment 8

```
[81]: ![ -e 'iris.data' ] || wget 'https://archive.ics.uci.edu/ml/
      ↪machine-learning-databases/iris/iris.data'
      ![ -e 'zip.train' ] || ( wget https://web.stanford.edu/~hastie/ElemStatLearn/
      ↪datasets/zip.train.gz && gzip -d zip.train.gz )
      ![ -e 'zip.test' ] || ( wget https://web.stanford.edu/~hastie/ElemStatLearn/
      ↪datasets/zip.test.gz && gzip -d zip.test.gz )
```

--2021-01-18 03:10:49--

https://web.stanford.edu/~hastie/ElemStatLearn/datasets/zip.train.gz

Resolving web.stanford.edu (web.stanford.edu)... 171.67.215.200,

2607:f6d0:0:925a::ab43:d7c8

Connecting to web.stanford.edu (web.stanford.edu)|171.67.215.200|:443...
connected.

HTTP request sent, awaiting response... 200 OK

Length: 1829071 (1,7M) [application/x-gzip]

Saving to: 'zip.train.gz'

zip.train.gz 100%[=====>] 1,74M 294KB/s in 8,6s

2021-01-18 03:10:58 (208 KB/s) - 'zip.train.gz' saved [1829071/1829071]

--2021-01-18 03:10:59--

https://web.stanford.edu/~hastie/ElemStatLearn/datasets/zip.test.gz

Resolving web.stanford.edu (web.stanford.edu)... 171.67.215.200,

2607:f6d0:0:925a::ab43:d7c8

Connecting to web.stanford.edu (web.stanford.edu)|171.67.215.200|:443...
connected.

HTTP request sent, awaiting response... 200 OK

Length: 439208 (429K) [application/x-gzip]

Saving to: 'zip.test.gz'

zip.test.gz 100%[=====>] 428,91K 190KB/s in 2,3s

2021-01-18 03:11:02 (190 KB/s) - 'zip.test.gz' saved [439208/439208]

```
[3]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.utils import shuffle
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[4]: class Classifier:

    def accuracy(self, labels, predictions):
        return np.mean(labels == predictions)

data = pd.read_csv("iris.data", header=None)
data.head(n=5)
```

```
[4]:      0      1      2      3      4
0  5.1  3.5  1.4  0.2  Iris-setosa
1  4.9  3.0  1.4  0.2  Iris-setosa
2  4.7  3.2  1.3  0.2  Iris-setosa
3  4.6  3.1  1.5  0.2  Iris-setosa
4  5.0  3.6  1.4  0.2  Iris-setosa
```

Splitting the data into training/test and according to their class memberships

```
[5]: X_train, X_test, y_train, y_test = train_test_split(data[list(range(4))],
↳data[4], test_size=0.2, random_state=None, stratify=data[4])

X_train_setosa = X_train[y_train=='Iris-setosa'].to_numpy()
X_train_versicolor = X_train[y_train=='Iris-versicolor'].to_numpy()
X_train_virginica = X_train[y_train=='Iris-virginica'].to_numpy()

X_test_setosa_v_v = X_test.to_numpy()
y_test_setosa_v_v = (y_test == 'Iris-setosa').astype(int).to_numpy()

X_test_versicolor_virginica = X_test[y_test!='Iris-setosa'].to_numpy()
y_test_versicolor_virginica = (y_test[y_test!='Iris-setosa'] ==
↳'Iris-versicolor').astype(int).to_numpy()
```

```
[74]: y_train1 = (y_train == 'Iris-setosa').astype(int).to_numpy()
X_train2 = X_train.to_numpy()[~y_train1.astype(bool)]
y_train2 = (y_train[~y_train1.astype(bool)] == 'Iris-versicolor').astype(int).
↳to_numpy()
```

```
[7]: class PerceptronClassifier(Classifier):
    def __init__(self, learning_rate=0.5, max_iters=100):
```

```

        self.learning_rate = learning_rate
        self.max_iters = max_iters
        self.errors = []
        self.weights = None

    def fit(self, X, y):
        X = np.hstack([X, np.ones((X.shape[0], 1))])
        self.weights = np.zeros(X.shape[1])
        for i in range(self.max_iters):
            error = 0
            for xj, yj in zip(X, y):
                dw = self.learning_rate * (yj - self.predict(xj[:-1].reshape(1,
↪-1)))

                self.weights += dw * xj
                error += int(dw != 0)
            self.errors.append(error)
            if error == 0:
                break
        return self

    def predict_z(self, X):
        return X @ self.weights[:-1] + self.weights[-1]

    def predict(self, X):
        return (self.predict_z(X) >= 0.0).astype(int)

```

```

[8]: %%time
model = PerceptronClassifier()
model.fit(X_train, y_train1)
y_hat_train1 = model.predict(X_train.to_numpy())
y_hat_test1 = model.predict(X_test_setosa_v_v)

```

CPU times: user 7.55 ms, sys: 523 µs, total: 8.07 ms
Wall time: 6.32 ms

```

[9]: print(f"Train accuracy: {model.accuracy(y_train1, y_hat_train1)}")
print(f"Test accuracy: {model.accuracy(y_test_setosa_v_v, y_hat_test1)}")

```

Train accuracy: 1.0
Test accuracy: 1.0

```

[63]: plot_axis = [1, 2]
cmap = 'plasma'
X_center = np.mean(X_train.to_numpy(), axis=0)
X_min = np.min(X_train.to_numpy(), axis=0)
X_max = np.max(X_train.to_numpy(), axis=0)

```

```

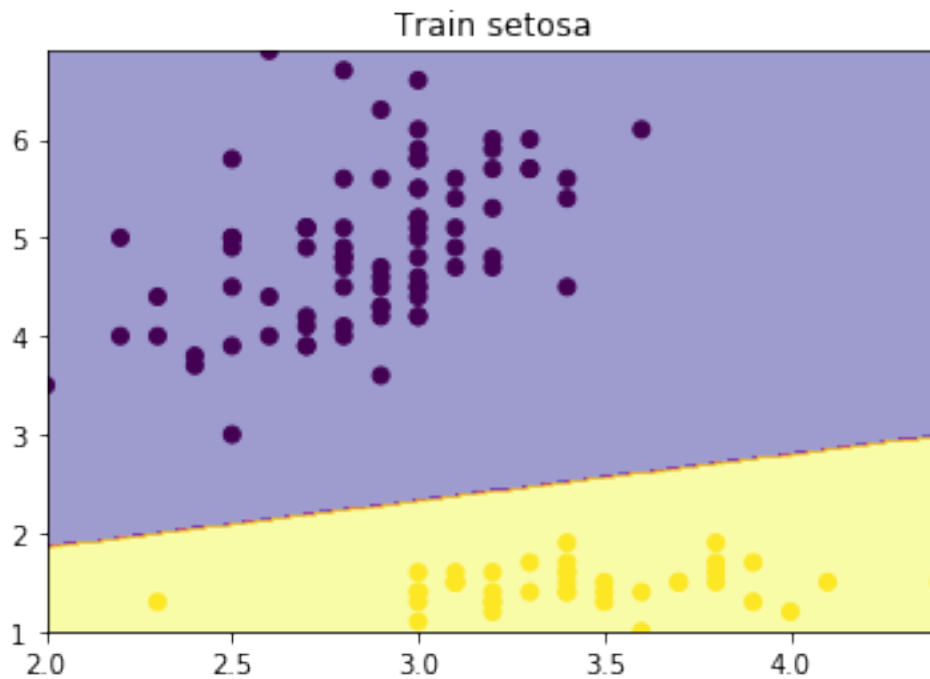
X_mesh = np.meshgrid(*([np.arange(X_min[i], X_max[i], 0.02) if i in plot_axis_
    ↳ else X_center[i] for i in range(len(X_min))])), indexing='ij')
X_mesh = [i.reshape((i.shape[plot_axis[0]], i.shape[plot_axis[1]])) for i in_
    ↳ X_mesh]
y_mesh1 = model.predict(np.array([X_mesh[i].ravel() for i in_
    ↳ range(len(X_mesh))]).T).reshape(X_mesh[0].shape)

```

```

[67]: plt.xlim(X_min[plot_axis[0]], X_max[plot_axis[0]])
plt.ylim(X_min[plot_axis[1]], X_max[plot_axis[1]])
plt.title('Train setosa')
plt.contourf(X_mesh[plot_axis[0]], X_mesh[plot_axis[1]], y_mesh1, alpha=.4,
    ↳ cmap=cmap, extend='both')
plt.scatter(X_train.to_numpy()[:, plot_axis[0]], X_train.to_numpy()[:,
    ↳ plot_axis[1]], c=y_train1)
plt.show()

```



```

[62]: %%time
model = PerceptronClassifier()
model.fit(X_train, y_train1)
y_hat_train1 = model.predict(X_train.to_numpy())
y_hat_test1 = model.predict(X_test_setosa_v_v)

```

CPU times: user 6.72 ms, sys: 52 μ s, total: 6.78 ms
 Wall time: 4.89 ms

```
[66]: plt.title('Test setosa')
plt.contourf(X_mesh[plot_axis[0]], X_mesh[plot_axis[1]], y_mesh1, alpha=.4,
             cmap=cmap, extend='both')
plt.scatter(X_test_setosa_v_v[:, plot_axis[0]], X_test_setosa_v_v[:,
             plot_axis[1]], c=y_test_setosa_v_v)
plt.show()
```



```
[75]: %%time
model2 = PerceptronClassifier()
model2.fit(X_train2, y_train2)
y_hat_train2 = model.predict(X_train2)
y_hat_test2 = model.predict(X_test_versicolor_virginica)
```

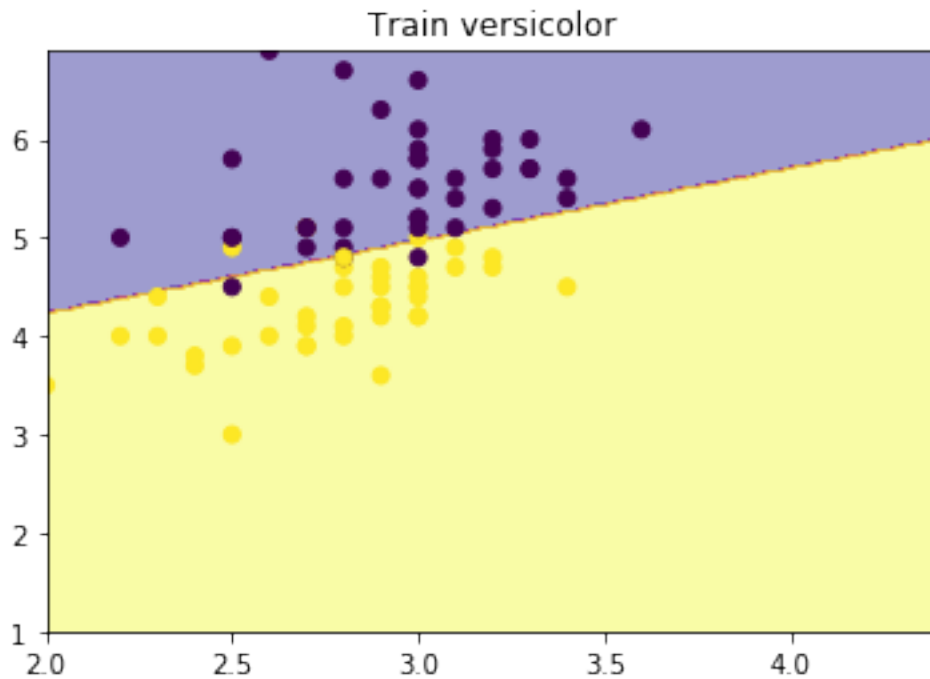
CPU times: user 161 ms, sys: 25.4 ms, total: 187 ms
Wall time: 144 ms

```
[76]: print(f"Train accuracy: {model.accuracy(y_train2, y_hat_train2)}")
print(f"Test accuracy: {model.accuracy(y_test_versicolor_virginica,
             y_hat_test2)}")
```

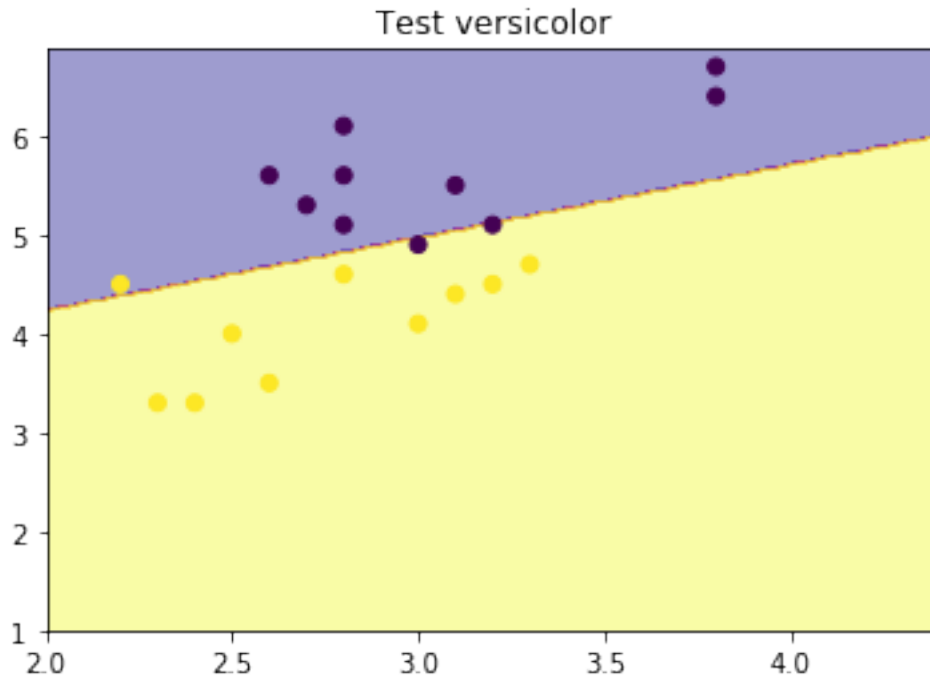
Train accuracy: 0.5
Test accuracy: 0.5

```
[77]: y_mesh2 = model2.predict(np.array([X_mesh[i].ravel() for i in
             range(len(X_mesh))]).T).reshape(X_mesh[0].shape)
```

```
[80]: plt.xlim(X_min[plot_axis[0]], X_max[plot_axis[0]])
plt.ylim(X_min[plot_axis[1]], X_max[plot_axis[1]])
plt.title('Train versicolor')
plt.contourf(X_mesh[plot_axis[0]], X_mesh[plot_axis[1]], y_mesh2, alpha=.4,
↪ cmap=cmap, extend='both')
plt.scatter(X_train2[:, plot_axis[0]], X_train2[:, plot_axis[1]], c=y_train2)
plt.show()
```



```
[79]: plt.title('Test versicolor')
plt.contourf(X_mesh[plot_axis[0]], X_mesh[plot_axis[1]], y_mesh2, alpha=.4,
↪ cmap=cmap, extend='both')
plt.scatter(X_test_versicolor_virginica[:, plot_axis[0]],
↪ X_test_versicolor_virginica[:, plot_axis[1]], c=y_test_versicolor_virginica)
plt.show()
```



Train dataset for the second model is not linearly separable, so `fit()` would never stop, if there was no iteration count limit.

```
[82]: training_data_zip = np.array(pd.read_csv('zip.train', sep=' ', header=None))
      test_data_zip = np.array(pd.read_csv('zip.test', sep=' ', header=None))

      X_zip_train, y_zip_train = training_data_zip[:,1:-1], training_data_zip[:,0]
      X_zip_test, y_zip_test = test_data_zip[:,1:], test_data_zip[:,0]
```

```
[287]: class MultilayerPerceptronClassifier(Classifier):
      def __init__(self, hidden_neurons_numbers, max_iters=50, learning_rate=0.3):
          self.hidden_neurons_numbers = tuple(hidden_neurons_numbers)
          self.max_iters = max_iters
          self.learning_rate = learning_rate

      def gradient(self, X, y, W):
          epsilon = 0.7
          L0 = self.loss(X, y, W)
          for i, w in enumerate(W):
              dw = np.zeros_like(w)
              gr = np.zeros_like(w)
              for j in range(w.shape[0]):
                  dw[j] = epsilon
                  gr[j] = (self.loss(X, y, W[:i] + [w + dw] + W[i + 1:]) - L0) / epsilon
          epsilon
```

```

        dw[j] = 0.0
    yield gr

def loss(self, X, y, W):
    return ((y - self.predict_with(W, X)) ** 2).mean()

def predict_with(self, W, X):
    for w in W:
        X = self.activate(X @ w)
    return X

def activate(self, a):
    return 1.0 / (1.0 + np.exp(-a))

def fit(self, X, Y):
    Y = Y.reshape(-1, 1)
    X = np.hstack([X, np.ones((X.shape[0], 1))])
    neuron_numbers = (X.shape[1],) + self.hidden_neurons_numbers + (1,)
    self.weights = list((np.random.rand(neuron_numbers[i], neuron_numbers[i_
→+ 1]) - .5) for i in range(len(neuron_numbers) - 1))
    for w in self.weights:
        w[:, -1].fill(0)
        w[-1, -1] = 1
    for it in range(self.max_iters):
        stop = True
        nabla_W = tuple(self.gradient(X, Y, self.weights))
        acc_delta = tuple(map(np.zeros_like, self.weights))
        for i, w in enumerate(self.weights):
            delta = -self.learning_rate * nabla_W[i]
            if not np.allclose(delta, 0):
                stop = False
            acc_delta[i].__iadd__(delta)
        for w, dw in zip(self.weights, acc_delta):
            w.__iadd__(dw)
            w[-1, :].fill(0)
            w[-1, -1] = 1
        if stop:
            break
    return self

def predict(self, X):
    return np.round(self.predict_with(self.weights, np.hstack([X, np.
→ones((X.shape[0], 1))])).astype(int)

```

```

[290]: %%time
np.random.seed(42)
model3 = MultilayerPerceptronClassifier((32, 16, 8))

```



```
model3.fit(X_zip_train[:64], y_zip_train[:64] == 4)
print(model3.accuracy(y_zip_train[:64] == 4, model3.predict(X_zip_train[:64])))
```

0.84375

CPU times: user 7.62 s, sys: 0 ns, total: 7.62 s

Wall time: 7.64 s